



Statistical Signal Processing

Lecture 10

chapter 3: Optimal Filtering

Wiener filtering

- FIR Wiener filtering
 - iterative solution: steepest-descent algorithm

chapter 4: Adaptive Filtering

- LMS algorithm
- Normalized LMS (NLMS) algorithm
- tracking behavior of LMS and RLS
- optimal tracking via Kalman filtering

chapter 5: Sinusoids in Noise



RLS Algorithm

- LS: replace the statistical averages by a time averages:

$$\xi_k(H) = \sum_{i=1}^k (x_i - H^T Y_i)^2 + (H - H_0)^T R_0 (H - H_0) ,$$

where the second term with $R_0 = R_0^T > 0$ allows for a proper initialization of the algorithm (the first term alone has a singular Hessian $(= 2 \sum_{i=1}^k Y_i Y_i^T)$ for $k < N$).

- We can rewrite

$$\begin{aligned} \xi_k(H) &= H^T \left(\sum_{i=1}^k Y_i Y_i^T \right) H - 2H^T \left(\sum_{i=1}^k Y_i x_i \right) + \sum_{i=1}^k x_i^2 + (H - H_0)^T R_0 (H - H_0) \\ &= H^T \left(R_0 + \sum_{i=1}^k Y_i Y_i^T \right) H - 2H^T \left(R_0 H_0 + \sum_{i=1}^k Y_i x_i \right) + \sum_{i=1}^k x_i^2 + H_0^T R_0 H_0 \\ &= H^T R_k H - 2H^T P_k + \sum_{i=1}^k x_i^2 + H_0^T R_0 H_0 \end{aligned}$$

where

$$\begin{aligned} R_k &= R_0 + \sum_{i=1}^k Y_i Y_i^T &= R_{k-1} + Y_k Y_k^T \\ P_k &= R_0 H_0 + \sum_{i=1}^k Y_i x_i &= P_{k-1} + Y_k x_k . \end{aligned}$$



Recursive Least-Squares Algorithm (2)

- By putting the gradient of $\xi_k(H)$ equal to zero and noting that the Hessian $2R_k > 0$, we find that the LS filter H_k that minimizes the LS criterion solves the following normal equations

$$R_k H_k = P_k .$$

To solve this set of equations at each time instant k would take $\mathcal{O}(N^3)$ operations at each time instant. In what follows, we shall derive the Recursive LS algorithm, which allows us, using information obtained at time $k-1$, to obtain H_k with only $\mathcal{O}(N^2)$ operations.

- we can rewrite $P_k = P_{k-1} + Y_k x_k$ as

$$\begin{aligned} R_k H_k &= R_{k-1} H_{k-1} + Y_k x_k \\ &= (R_k - Y_k Y_k^T) H_{k-1} + Y_k x_k \\ &= R_k H_{k-1} + Y_k \epsilon_k^p \end{aligned}$$

where $\epsilon_k^p = x_k - H_{k-1}^T Y_k$ as in the LMS algorithm. This leads immediately to

$$H_k = H_{k-1} + R_k^{-1} Y_k \epsilon_k^p$$

where $R_k^{-1} Y_k$ is called the Kalman gain (the RLS algorithm is a special case of the so-called Kalman filter).



Recursive Least-Squares Algorithm (3)

- Clearly, the RLS algorithm requires the recursive update of R_k^{-1} . This can be obtained using the Matrix Inversion Lemma:

$$\begin{aligned} R_k^{-1} &= (R_{k-1} + Y_k Y_k^T)^{-1} \\ &= R_{k-1}^{-1} - R_{k-1}^{-1} Y_k (1 + Y_k^T R_{k-1}^{-1} Y_k)^{-1} Y_k^T R_{k-1}^{-1} . \end{aligned}$$

This equation allows us to obtain R_k^{-1} from R_{k-1}^{-1} and Y_k using $\mathcal{O}(N^2)$ operations. When multiplying both sides with Y_k to the right, we obtain

$$R_k^{-1} Y_k = R_{k-1}^{-1} Y_k (1 + Y_k^T R_{k-1}^{-1} Y_k)^{-1} .$$

We find for the *a posteriori* error

$$\epsilon_k = x_k - H_k^T Y_k = (1 - Y_k^T R_k^{-1} Y_k) \epsilon_k^p = (1 + Y_k^T R_{k-1}^{-1} Y_k)^{-1} \epsilon_k^p .$$

- All this can be formulated as the RLS algorithm:

$$\left\{ \begin{array}{l} \epsilon_k^p = x_k - H_{k-1}^T Y_k \\ \epsilon_k = \epsilon_k^p (1 + Y_k^T R_{k-1}^{-1} Y_k)^{-1} \\ H_k = H_{k-1} + R_{k-1}^{-1} Y_k \epsilon_k \\ R_k^{-1} = R_{k-1}^{-1} - R_{k-1}^{-1} Y_k (1 + Y_k^T R_{k-1}^{-1} Y_k)^{-1} Y_k^T R_{k-1}^{-1} . \end{array} \right.$$



Recursive Least-Squares Algorithm (4)

- The initial values for R_k^{-1} and H_k are R_0^{-1} and H_0 . Compared to the LMS algorithm, the scalar stepsize μ gets replaced by a matrix stepsize R_k^{-1} . The RLS algorithm takes $\mathcal{O}(N^2)$ operations while the LMS algorithm takes only $2N$ operations. However, it converges much faster.

- *performance analysis* : with $x_k = H^o Y_k + \tilde{x}_k$ (and $R_0 = 0$), we get

$$R_k H_k = P_k = \sum_{i=1}^k Y_i x_i = R_k H^o + \sum_{i=1}^k Y_i \tilde{x}_i . \text{ Hence}$$

$$\widetilde{H}_k = H^o - H_k = -R_k^{-1} \sum_{i=1}^k Y_i \tilde{x}_i$$

From this, we obtain

$$C_k \triangleq E \widetilde{H}_k \widetilde{H}_k^T = \sigma_{\tilde{x}}^2 R_k^{-1} .$$

Since R_k^{-1} behaves as $1/k$, we see that C_k converges to zero as $1/k$.

- *Exponential Weighting* In order to be able to track a possibly time-varying $H^o = H_k^o$, one introduces an exponential forgetting factor $\lambda \in (0, 1)$ into the cost function to obtain

$$\xi_k(H) = \sum_{i=1}^k \lambda^{k-i} (x_i - H^T Y_i)^2 + \lambda^k (H - H_0)^T R_0 (H - H_0) .$$

This implies that the past (and in particular the initial conditions H_0, R_0) is forgotten exponentially fast with a window with time constant $1/(1-\lambda)$.



Recursive Least-Squares Algorithm (5)

- Wiener filtering: x_k and y_k are two joint stochastic processes and we're trying to estimate x_k from the y_k using a LMMSE estimator. For an FIR Wiener filter, there are a finite set of coefficients H^o involved in this LMMSE estimator.
RLS approach: replaced statistical averages with temporal averages.

Parameter estimation interpretation

- Assume now our usual model for the *measurements* x_k ,

$$x_k = H^{oT} Y_k + \tilde{x}_k$$

where the \tilde{x}_k are iid with zero mean and variance $\sigma_{\tilde{x}}^2$. Consider here $\{y_k\}$ as a deterministic signal, so the only randomness comes from the $\{\tilde{x}_k\}$. The H^o are the unknown parameters governing the model.

- The analysis of the RLS algorithm is much simpler than that of the LMS algorithm since for each k the RLS solution H_k coincides with the solution of a Least-Squares problem with a closed-form solution: $H_k = R_k^{-1} P_k$.
- Assume now $k \geq N$, $H_0 = 0$, $R_0 = 0$, and that R_k is nonsingular. The performance of the least-squares estimate is simple to analyze and leads to

$$C_k \triangleq E \widetilde{H}_k \widetilde{H}_k^T = \sigma_{\tilde{x}}^2 R_k^{-1} .$$

- If \tilde{x}_k Gaussian, $\Rightarrow H_k = \text{ML estimate of } H^o$ (efficient, $C_k = \text{CRB}$).



Recursive Least-Squares Algorithm (6)

A Bayesian Context - A Priori Information

- Instead of treating the filter coefficients H^o as unknown constant parameters, we could also consider H^o as a stochastic parameter vector about which we have some prior information, possibly from previous adaptive filtering experience. Assume now that, prior to obtaining the measurements x_1, x_2, \dots , we know that H^o has a distribution with mean $E H^o = H_0$ and covariance $E (H^o - H_0) (H^o - H_0)^T = C_0$. So now the randomness in the x_k comes from both the \tilde{x}_k and H^o .
- The problem formulation can now be recognized to be one of a *Bayesian Linear Model*. The AMMSE estimator can be shown to be the filter estimate resulting from the original RLS criterion with $R_0 = \sigma_{\tilde{x}}^2 C_0^{-1}$. $C_k = E \widetilde{H}_k \widetilde{H}_k^T$ now satisfies

$$C_k^{-1} = \sigma_{\tilde{x}}^{-2} R_k = \sigma_{\tilde{x}}^{-2} \sum_{i=1}^k Y_i Y_i^T + C_0^{-1} .$$

Note that C_k^{-1} is an increasing function of C_0^{-1} and hence C_k is a decreasing function of C_0^{-1} and hence of R_0 .

- So we see that H_0 and R_0 in the LS cost function have the interpretation of the prior mean and the inverse of the prior covariance of H^o . We'll choose R_0 small if we don't have a lot of confidence in our prior guess H_0 (C_0 big). In practice, R_0 is often chosen as $R_0 = \eta I_N$.



Other Adaptive Filtering Algorithms

- Fast RLS algorithms: Fast Transversal Filter (FTF) algorithm ($8N$), Fast Lattice/QR Algorithms ($\mathcal{O}(N)$ complexity)
- LMS with prewhitened input
- block processing/frequency domain LMS
- subband structures
- Fast Newton Transversal Filter (FNTF): replace R^{-1} in RLS by a banded matrix (appropriate for AR processes, hence speech)
- projection algorithms (like NLMS) on an extended subspace of L input vectors (FAP: Fast Affine Projection: complexity $2N + \mathcal{O}(L^2)$ or $2N + \mathcal{O}(L)$)
- Fast Subsampled Updating (FSU) versions of LMS and FTF: introduce some delay to reduce complexity below $\mathcal{O}(N)$
- multistage Wiener filter / polynomial expansion:

$$r_0 R^{-1} = \left[\frac{1}{r_0} R \right]^{-1} = \left[\underbrace{I}_{\text{diagonal}} + \underbrace{\left(\frac{1}{r_0} R - I \right)}_{\text{off-diagonal part}} \right]^{-1} = \sum_{i=0}^{\infty} \left(I - \frac{1}{r_0} R \right)^i = \sum_{i=0}^{\infty} \alpha_i R^i$$

- convergence speed (RLS best) versus tracking speed (FAP best?)

Initial Convergence RLS

- Consider now $H_0 \neq 0, R_0 \neq 0$,

$$R_k = R_0 + R_{1:k}, R_{1:k} = Y_{1:k} Y_{1:k}^T, Y_{1:k} = [Y_1 \cdots Y_k], P_k = R_0 H_0 + P_{1:k}$$

- $\tilde{H}_k = H^o - H_k = H^o - R_k^{-1} P_k = (R_0 + R_{1:k})^{-1} (R_0 \tilde{H}_0 - \sum_{i=1}^k Y_i \tilde{x}_i)$

- $C_k = E \tilde{H}_k \tilde{H}_k^T$ hence

$$\begin{aligned} C_k &= (R_0 + R_{1:k})^{-1} R_0 \tilde{H}_0 \tilde{H}_0^T R_0 (R_0 + R_{1:k})^{-1} + \sigma_{\tilde{x}}^2 (R_0 + R_{1:k})^{-1} R_{1:k} (R_0 + R_{1:k})^{-1} \\ &= \underbrace{\sigma_{\tilde{x}}^2 (R_0 + R_{1:k})^{-1}}_{\sim \frac{1}{k}} + \underbrace{(R_0 + R_{1:k})^{-1} (R_0 \tilde{H}_0 \tilde{H}_0^T R_0 - \sigma_{\tilde{x}}^2 R_0) (R_0 + R_{1:k})^{-1}}_{\sim \frac{1}{k^2} \text{ due to initialization}} \end{aligned}$$

- noiseless case: $\sigma_{\tilde{x}}^2 = 0$

$$C_k = (R_0 + R_{1:k})^{-1} R_0 \tilde{H}_0 \tilde{H}_0^T R_0 (R_0 + R_{1:k})^{-1}$$

- initial convergence: $1 \leq k < N$, consider $R_0 = \eta I$

$$C_k = \eta^2 (\eta I + R_{1:k})^{-1} \tilde{H}_0 \tilde{H}_0^T (\eta I + R_{1:k})^{-1}$$

Initial Convergence RLS (2)

- *Singular Value Decomposition (SVD)*: $Y_{1:k}$, $N \times k$ assumed full column rank

$$Y_{1:k} = V \Sigma U^T \quad V^T V = I_k, U^{-1} = U^T, \Sigma = \text{diag}\{\sigma_1, \dots, \sigma_k\}$$

$\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_k \geq 0$ “singular values” full column rank $\leftrightarrow \sigma_k > 0$

- *Moore-Penrose pseudo-inverse*: $Y_{1:k}^+ = U \Sigma^{-1} V^T = (Y_{1:k}^T Y_{1:k})^{-1} Y_{1:k}^T$
- projection on column space: $P_{Y_{1:k}} = Y_{1:k} Y_{1:k}^+ = V V^T$
- $V^+ = V^T$ $P_{Y_{1:k}} = V V^T = V V^+ = P_V$ $P_V^+ = P_V$
- eigendecomposition: $R_{1:k} = Y_{1:k} Y_{1:k}^T = V \Sigma^2 V^T$
- let V^\perp be such that $[V \ V^\perp]$ is orthogonal:

$$[V \ V^\perp][V \ V^\perp]^T = I = V V^T + V^\perp V^{\perp T} = P_V + P_{V^\perp} = P_V + P_V^\perp$$

$P_V^\perp = I - P_V = P_{V^\perp}$, V^\perp spans orthogonal complement of V

- SVD alternatively: $Y_{1:k} = [V \ V^\perp] \begin{bmatrix} \Sigma \\ 0 \end{bmatrix} U^T$, $\begin{bmatrix} \Sigma \\ 0 \end{bmatrix}^+ = [\Sigma^+ \ 0]$, $\sigma^+ = \begin{cases} 1/\sigma & , \sigma > 0 \\ 0 & , \sigma = 0 \end{cases}$
- eigendecomposition projection:

$$P_V = V I V^T = [V \ V^\perp] \begin{bmatrix} I & 0 \\ 0 & 0 \end{bmatrix} [V \ V^\perp]^T \quad \text{eigenvalues are 1 or 0}$$

Initial Convergence RLS (3)

- let $\eta \ll \sigma_k^2$ be small, then

$$\begin{aligned}
 \eta(\eta I + R_{1:k})^{-1} &= \eta(\eta V^\perp V^{\perp T} + \eta V V^T + V \Sigma^2 V^T)^{-1} \approx \eta(\eta V^\perp V^{\perp T} + V \Sigma^2 V^T)^{-1} \\
 &= \eta \left([V \ V^\perp] \begin{bmatrix} \Sigma^2 & 0 \\ 0 & \eta I \end{bmatrix} [V \ V^\perp]^T \right)^{-1} = \eta [V \ V^\perp] \begin{bmatrix} \Sigma^2 & 0 \\ 0 & \eta I \end{bmatrix}^{-1} [V \ V^\perp]^T \\
 &= \eta V \Sigma^{-2} V^T + V^\perp V^{\perp T} = \eta R_{1:k}^+ + \mathbf{P}_{R_{1:k}}^\perp
 \end{aligned}$$

where $R_{1:k}^+ = Y_{1:k} (Y_{1:k}^T Y_{1:k})^{-1} Y_{1:k}^T$ and $\mathbf{P}_{R_{1:k}} = \mathbf{P}_{Y_{1:k}}$

- hence

$$C_k = (\eta R_{1:k}^+ + \mathbf{P}_{R_{1:k}}^\perp) \widetilde{H}_0 \widetilde{H}_0^T (\eta R_{1:k}^+ + \mathbf{P}_{R_{1:k}}^\perp)$$

$\eta R_{1:k}^+ \widetilde{H}_0$ reduced to $\mathcal{O}(\eta)$ in k -dim. subspace, column space of $Y_{1:k}$

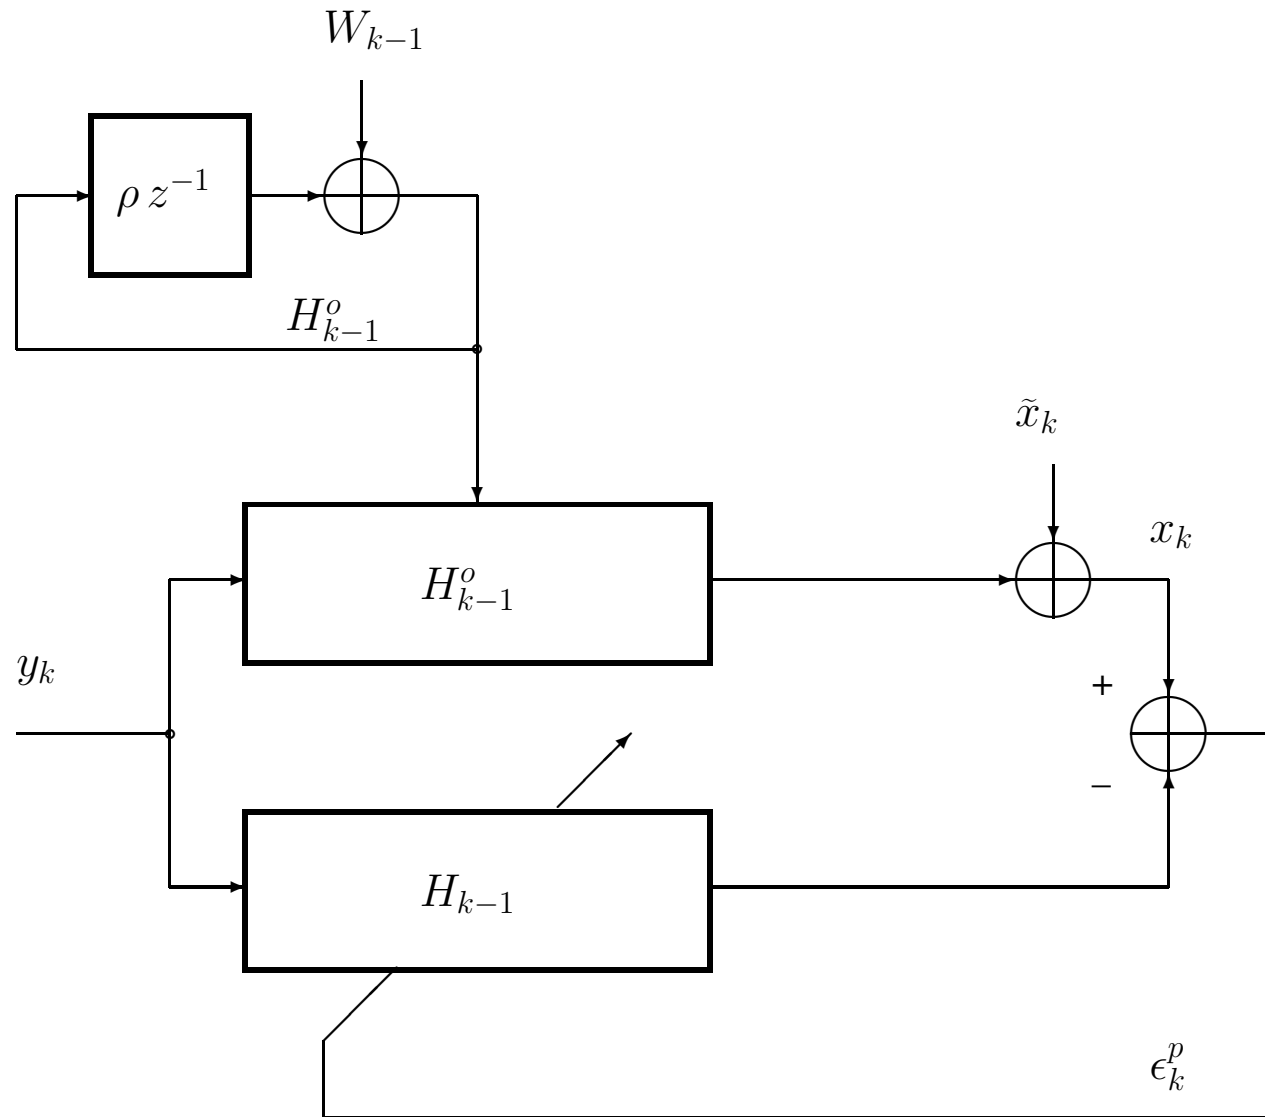
$\mathbf{P}_{R_{1:k}}^\perp \widetilde{H}_0$ unchanged in $(N - k)$ -dim. orthogonal complement

- C_k rank 1 (noiseless case): only the mean of \widetilde{H}_k needs to converge
- RLS: the mean of \widetilde{H}_k has essentially converged (filter estimate unbiased) after $k = N \Rightarrow$ very fast (mean dominates initial convergence in general)

LMS: the mean needs to converge exponentially, dynamics of steepest-descent



Tracking Time-Varying Filters



Time-Varying System Identification Set-Up

- system processes:

$$\begin{aligned} x_k &= Y_k^T H_{k-1}^o + \tilde{x}_k & E \tilde{x}_k \tilde{x}_i &= \xi^o \delta_{ki} \\ H_k^o &= \rho H_{k-1}^o + W_k & E W_k W_i^T &= Q \delta_{ki} \\ \tilde{H}_k &= H_k^o - H_k & E W_k \tilde{x}_i &= 0 \end{aligned}$$

- time-varying filter modeled as AR(1) process, requires $|\rho| < 1$ for stationarity
→ stationary case of nonstationarity
- adaptive filter a priori error signal:

$$\epsilon_k^p = x_k - Y_k^T H_{k-1} = Y_k^T \tilde{H}_{k-1} + \tilde{x}_k$$

- learning curve: (independence assumption)

$$\xi_k = E(\epsilon_k^p)^2 = \xi^o + \xi_k^e = \xi^o(1 + \mathcal{M}) , \quad \xi_k^e = \text{tr}\{R_{YY} C_{k-1}\} , \quad C_k = E \tilde{H}_k \tilde{H}_k^T$$

- consider $\frac{1}{1-\rho} \gg$ adaptation time constants so that we can take $\rho = 1$ for the analysis

Tracking Analysis LMS

- filter deviation recursion:

$$\begin{aligned}\widetilde{H}_k &= \widetilde{H}_{k-1} - \mu \epsilon_k^p Y_k + W_k \\ &= (I - \mu Y_k Y_k^T) \widetilde{H}_{k-1} - \mu \widetilde{x}_k Y_k + W_k \\ &\approx (I - \mu R_{YY}) \widetilde{H}_{k-1} - \mu \widetilde{x}_k Y_k + W_k\end{aligned}$$

where we introduced the averaging approach in the last step

- filter error correlation matrix recursion:

$$C_k = (I - \mu R_{YY}) C_{k-1} (I - \mu R_{YY}) + \mu^2 \xi^o R_{YY} + Q$$

- the stationary nonstationarity combined with a constant stepsize leads to a steady-state, for which we get (with small μ):

$$R_{YY} C_\infty + C_\infty R_{YY} = \mu \xi^o R_{YY} + \frac{1}{\mu} Q$$

- steady-state misadjustment: $\mathcal{M}_{LMS} = \underbrace{\frac{\mu}{2} \text{tr} R_{YY}}_{\text{estimation noise}} + \underbrace{\frac{1}{2\mu \xi^o} \text{tr} Q}_{\text{lag noise}}$

Tracking Analysis RLS

- filter deviation recursion:

$\lambda < 1$ to allow tracking

$$\begin{aligned}\widetilde{H}_k &= \widetilde{H}_{k-1} - R_k^{-1} Y_k \epsilon_k^p + W_k \\ &= (I - R_k^{-1} Y_k Y_k^T) \widetilde{H}_{k-1} - R_k^{-1} Y_k \widetilde{x}_k + W_k\end{aligned}$$

- after averaging in steady-state, assuming small $1 - \lambda$

$$(I - R_k^{-1} Y_k Y_k^T = \lambda R_k^{-1} R_{k-1} \approx \lambda I, \quad R_k^{-1} \approx (1 - \lambda) R_{YY}^{-1})$$

$$\widetilde{H}_k = \lambda \widetilde{H}_{k-1} - (1 - \lambda) R_{YY}^{-1} Y_k \widetilde{x}_k + W_k \quad (\text{dynamics indep. of } R_{YY})$$

- filter error correlation matrix recursion:

$$C_k = \lambda^2 C_{k-1} + (1 - \lambda)^2 \xi^o R_{YY}^{-1} + Q$$

- which leads to the steady-state value (assuming small $1 - \lambda$)

$$C_\infty = \frac{1 - \lambda}{2} \xi^o R_{YY}^{-1} + \frac{1}{2(1 - \lambda)} Q$$

- steady-state misadj.: $\mathcal{M}_{RLS} = \underbrace{\frac{1 - \lambda}{2} N}_{\text{estimation noise}} + \underbrace{\frac{1}{2(1 - \lambda) \xi^o} \text{tr}\{R_{YY} Q\}}_{\text{lag noise}}$

Tracking Optimization & LMS-RLS Comparison

- stepsize μ , $1 - \lambda$ design result of compromise between:
 estimation noise: finite stepsize prevents convergence, consistency
 lag noise: small stepsize leads to lowpass filtering and to a filter estimate that lags behind the true filter

- LMS: $\mu^{opt} = \sqrt{\frac{\text{tr}Q}{\xi^o \text{tr}R_{YY}}}$, $\mathcal{M}_{LMS}^{opt} = \sqrt{\frac{\text{tr}R_{YY} \text{tr}Q}{\xi^o}}$

- RLS: $\lambda^{opt} = 1 - \sqrt{\frac{\text{tr}\{R_{YY}Q\}}{N \xi^o}}$, $\mathcal{M}_{RLS}^{opt} = \sqrt{\frac{N \text{tr}\{R_{YY}Q\}}{\xi^o}}$

- comparison:

$$\frac{\mathcal{M}_{LMS}^{opt}}{\mathcal{M}_{RLS}^{opt}} = \sqrt{\frac{\text{tr}R_{YY} \text{tr}Q}{N \text{tr}\{R_{YY}Q\}}}$$

$$Q = \begin{cases} q I & : \text{equal performance, at least for small } q \\ q R_{YY} & : \text{LMS is better} \\ q R_{YY}^{-1} & : \text{RLS is better} \end{cases}$$

- faster initial convergence of RLS could be exploited for *jumping* parameters, if window size properly adapted



Optimal Tracking via Kalman Filtering

- *state-space model*: state = AR(1) vector process

$$\begin{array}{l} \text{state equation} \quad H_k^o = A_k H_{k-1}^o + W_k \\ \text{measurement equation} \quad x_k = Y_k^T H_{k-1}^o + \tilde{x}_k \end{array}, \quad E \begin{bmatrix} W_k \\ \tilde{x}_k \end{bmatrix} \begin{bmatrix} W_i \\ \tilde{x}_i \end{bmatrix}^T = \begin{bmatrix} Q_k & P_k^T \\ P_k & \xi_k^o \end{bmatrix} \delta_{ki}$$

time-varying (at the very least due to Y_k), usually $P_k = 0$

state noise: W_k , measurement noise: \tilde{x}_k , state transition matrix A_k

- Kalman filter (KF): estimates recursively in time the *state* H_{k-1}^o on the basis of the *measurements* x_0, \dots, x_k in a LMMSE sense.
Sources of randomness: W_k, \tilde{x}_k . Signal y_k treated as deterministic.
- special case: $H_k^o = H_{k-1}^o \Rightarrow$ Kalman filter \rightarrow RLS algorithm
- RLS with exponential weighting can be interpreted as KF for the case of some non-zero Q_k
- in the time-invariant case (x_k and H_k^o jointly stationary apart from initial conditions), the KF converges to the causal Wiener filter.